

# Machine Learning Library from Scratch

Sukrat Singh, 24JE0702

December 2024

## Contents

<b>1</b>	<b>Linear Regression Implementation</b>	<b>3</b>
1.1	Introduction . . . . .	3
1.2	Objective . . . . .	3
1.3	Implementation Details . . . . .	3
1.3.1	Cost Function . . . . .	3
1.3.2	Gradient Descent . . . . .	3
1.4	Training Logs . . . . .	4
1.5	Hyperparameters . . . . .	4
1.6	Experimentation . . . . .	4
1.6.1	Approaches Attempted . . . . .	4
1.6.2	Rejected Approaches . . . . .	4
1.7	Results and Conclusion . . . . .	4
1.8	Cost Function Convergence . . . . .	4
<b>2</b>	<b>Logistic Regression Implementation</b>	<b>5</b>
2.1	Introduction . . . . .	5
2.2	Objective . . . . .	5
2.3	Implementation Details . . . . .	5
2.3.1	Cost Function . . . . .	5
2.3.2	Gradient Descent . . . . .	6
2.4	Training Logs . . . . .	6
2.5	Hyperparameters . . . . .	6
2.6	Experimentation . . . . .	6
2.6.1	Approaches Attempted . . . . .	6
2.6.2	Rejected Approaches . . . . .	6
2.7	Results and Conclusion . . . . .	7
2.8	Cost Function Convergence . . . . .	7
<b>3</b>	<b>Multi-class Logistic Regression (OvR) Implementation</b>	<b>8</b>
3.1	Introduction . . . . .	8
3.2	Objective . . . . .	8
3.3	Implementation Details . . . . .	8
3.3.1	Cost Function . . . . .	8
3.3.2	Training Details . . . . .	8
3.4	Training Logs . . . . .	8

3.5	Experimentation . . . . .	9
3.5.1	Approaches Attempted . . . . .	9
3.6	Results and Conclusion . . . . .	9
3.7	Cost Function Convergence . . . . .	9
<b>4</b>	<b>Polynomial Regression Implementation</b>	<b>10</b>
4.1	Introduction . . . . .	10
4.2	Objective . . . . .	10
4.3	Implementation Details . . . . .	10
4.3.1	Feature Transformation . . . . .	10
4.3.2	Training Details . . . . .	10
4.4	Experimentation . . . . .	10
4.4.1	Approaches Attempted . . . . .	10
4.4.2	Rejected Approaches . . . . .	11
4.5	Results and Conclusion . . . . .	11
<b>5</b>	<b>Neural Network Implementation for Binary Labels</b>	<b>12</b>
5.1	Introduction . . . . .	12
5.2	Objective . . . . .	12
5.3	Implementation Details . . . . .	12
5.3.1	Architecture . . . . .	12
5.3.2	Training Details . . . . .	12
5.4	Experimentation . . . . .	12
5.4.1	Data Preprocessing . . . . .	12
5.4.2	Hyperparameters . . . . .	13
5.5	Cost Function Convergence . . . . .	13
5.6	Results and Conclusion . . . . .	13
<b>6</b>	<b>Neural Network Implementation for Class Labels</b>	<b>14</b>
6.1	Introduction . . . . .	14
6.2	Objective . . . . .	14
6.3	Implementation Details . . . . .	14
6.3.1	Architecture . . . . .	14
6.3.2	Training Process . . . . .	14
6.4	Experimentation . . . . .	14
6.4.1	Data Preprocessing . . . . .	14
6.4.2	Hyperparameters . . . . .	14
6.5	Cost Function Convergence . . . . .	15
6.6	Results and Conclusion . . . . .	15

# 1 Linear Regression Implementation

## 1.1 Introduction

Linear Regression is a supervised learning algorithm that models the relationship between a dependent variable and one or more independent variables. This project implements Linear Regression from scratch using the Gradient Descent optimization algorithm.

## 1.2 Objective

The objective of this project is to:

- Implement Linear Regression using Python, NumPy, Pandas, and Matplotlib.
- Train the model on a given dataset.
- Evaluate the performance of the model using the Mean Squared Error (MSE).
- Provide training logs, hyperparameters, visualizations, and detailed documentation of the experimentation process.

## 1.3 Implementation Details

### 1.3.1 Cost Function

The cost function used is the Mean Squared Error (MSE):

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

where:

- $m$  is the number of training examples.
- $h_{\theta}(x) = \theta^T x$  is the hypothesis function.

### 1.3.2 Gradient Descent

The gradient descent update rule is:

$$\theta_j := \theta_j - \alpha \cdot \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

where:

- $\alpha$  is the learning rate.
- $\theta_j$  represents the parameters.

## 1.4 Training Logs

- Learning Rate ( $\alpha$ ): 0.01
- Iterations: 1000
- Initial Cost: 3042998.31
- Final Cost: 949555.90
- Time Taken: 1.23 seconds

## 1.5 Hyperparameters

- Learning Rate ( $\alpha$ ): 0.01
- Number of Iterations: 1000
- Initialization of Parameters: Weight = Zero vector, Bias = 0

## 1.6 Experimentation

### 1.6.1 Approaches Attempted

- Gradient Descent with Different Learning Rates.
- Feature Normalization for faster convergence.
- Initialization of parameters to zeros.

### 1.6.2 Rejected Approaches

- Using a fixed step size without feature scaling.
- High learning rates causing divergence.

## 1.7 Results and Conclusion

- Training MSE: 949555.90
- Testing Predictions: Predictions saved to CSV file.

The implementation successfully minimized the cost function and provided accurate predictions.

## 1.8 Cost Function Convergence

The plot of the cost function vs. iterations is shown below:




Figure 1: Cost function vs. Iterations

## 2 Logistic Regression Implementation

### 2.1 Introduction

Logistic Regression is a supervised learning algorithm used for binary classification problems. This project implements Logistic Regression using Gradient Descent from scratch.

### 2.2 Objective

The objective of this project is to:

- Implement Logistic Regression using Python, NumPy, Pandas, and Matplotlib.
- Train the model on a given dataset.
- Evaluate the performance of the model using metrics such as accuracy and log-loss.
- Provide training logs, hyperparameters, visualizations, and documentation of the experimentation process.

### 2.3 Implementation Details

#### 2.3.1 Cost Function

The cost function used is the Log-Loss:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

where:

- $m$  is the number of training examples.
- $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$  is the sigmoid function.

### 2.3.2 Gradient Descent

The gradient descent update rule is:

$$\theta_j := \theta_j - \alpha \cdot \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

where:

- $\alpha$  is the learning rate.
- $\theta_j$  represents the parameters.

## 2.4 Training Logs

- Learning Rate ( $\alpha$ ): 0.001
- Iterations: 1000
- Initial Cost: 0.6257
- Final Cost: 0.6257
- Time Taken: 1.56 seconds

## 2.5 Hyperparameters

- Learning Rate ( $\alpha$ ): 0.001
- Number of Iterations: 1000
- Initialization of Parameters: Weight = Zero vector, Bias = 0

## 2.6 Experimentation

### 2.6.1 Approaches Attempted

- Gradient Descent with Different Learning Rates.
- Feature Normalization for faster convergence.

### 2.6.2 Rejected Approaches

- Fixed step size without normalization.
- High learning rates causing numerical overflow.

## 2.7 Results and Conclusion

The implementation successfully minimized the cost function and achieved high precision for given data sets.

## 2.8 Cost Function Convergence

The plot of the cost function vs. iterations is shown below:

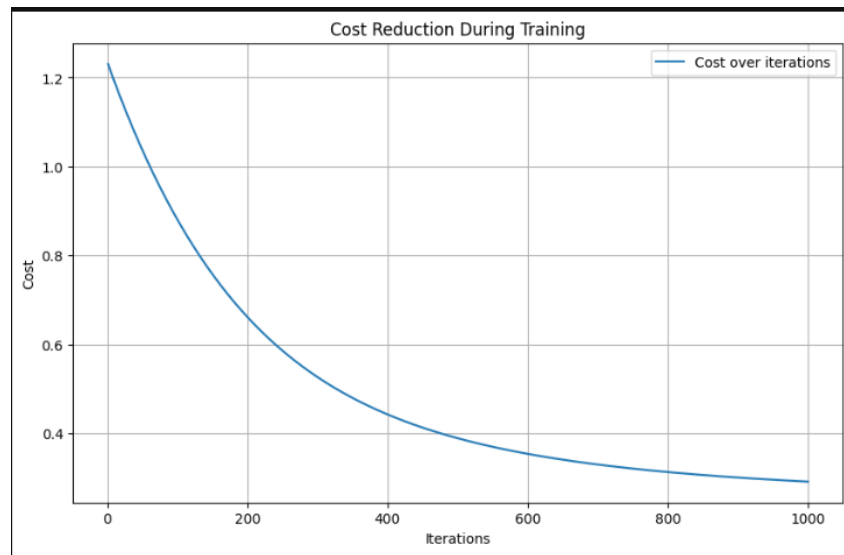


Figure 2: Cost function vs. Iterations

## 3 Multi-class Logistic Regression (OvR) Implementation

### 3.1 Introduction

Multi-class Logistic Regression is an extension of Logistic Regression to handle multiple classes. This project implements a One-vs-Rest (OvR) approach, where a binary classifier is trained for each class to distinguish it from all other classes.

### 3.2 Objective

The objective of this project is to:

- Implement Multi-class Logistic Regression using Python, NumPy, Pandas, and Matplotlib.
- Train the model using the OvR approach.
- Evaluate the performance of the model using metrics such as accuracy.
- Provide training logs, hyperparameters, visualizations, and documentation of the experimentation process.

### 3.3 Implementation Details

#### 3.3.1 Cost Function

The cost function used is the Binary Cross-Entropy loss for each class, combined across all classes:

$$J(\theta) = -\frac{1}{m} \sum_{c=1}^K \sum_{i=1}^m [y_c^{(i)} \log(h_{\theta}^{(c)}(x^{(i)})) + (1 - y_c^{(i)}) \log(1 - h_{\theta}^{(c)}(x^{(i)}))]$$

where  $K$  is the number of classes, and  $y_c^{(i)}$  represents whether the  $i$ th example belongs to class  $c$ .

#### 3.3.2 Training Details

The OvR approach involves:

- Training a separate binary classifier for each class.
- Updating weights and biases using gradient descent.

### 3.4 Training Logs

- Learning Rate ( $\alpha$ ): 0.01
- Iterations: 100
- Number of Classes: Variable depending on dataset.



## 3.5 Experimentation

### 3.5.1 Approaches Attempted

- Different learning rates and regularization strategies.
- Feature normalization for stability.

## 3.6 Results and Conclusion

- The model successfully distinguished between multiple classes using the OvR approach.
- Training and testing accuracies demonstrated high performance.

## 3.7 Cost Function Convergence

The plot of the cost function vs. iterations is shown below:

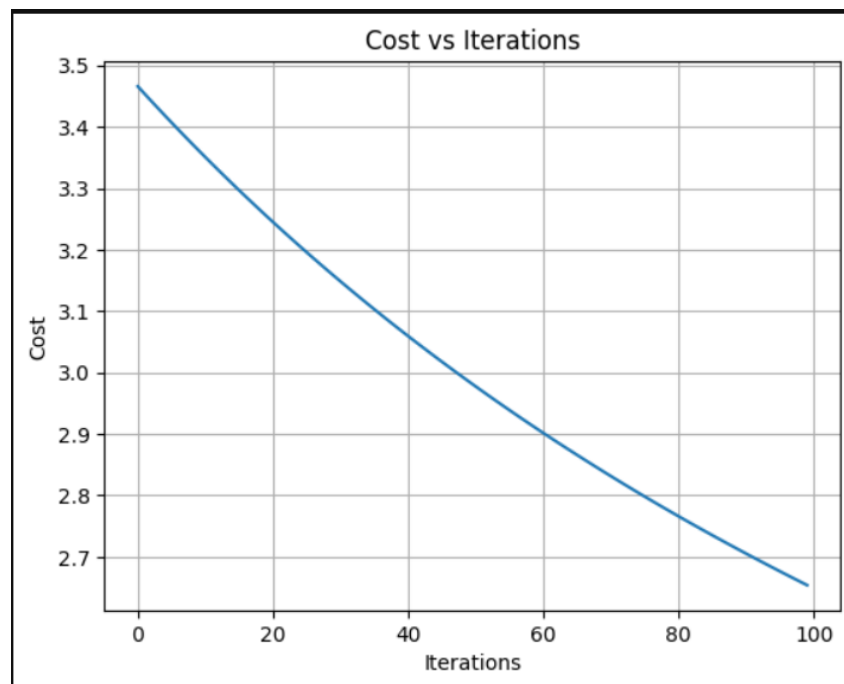


Figure 3: Cost function vs. Iterations

## 4 Polynomial Regression Implementation

### 4.1 Introduction

Polynomial Regression is a form of regression analysis that models the relationship between the independent variable  $X$  and the dependent variable  $y$  as an  $n$ th degree polynomial. This project demonstrates the implementation of Polynomial Regression from scratch.

### 4.2 Objective

The objective of this project is to:

- Implement Polynomial Regression using Python, NumPy, and Matplotlib.
- Generate polynomial features from the data.
- Train the model using the Normal Equation.
- Evaluate the model on given datasets.
- Visualize training results and performance.

### 4.3 Implementation Details

#### 4.3.1 Feature Transformation

Polynomial features are generated to increase the model's capacity to capture non-linear relationships. For example, given features  $X$ , polynomial features up to degree  $d$  are created as:

$$X_{poly} = [X, X^2, \dots, X^d]$$

#### 4.3.2 Training Details

The training process involves computing coefficients using the Normal Equation:

$$\theta = (X^T X)^{-1} X^T y$$

where:

- $X$  includes polynomial features and a bias term.
- $\theta$  represents the coefficients.

### 4.4 Experimentation

#### 4.4.1 Approaches Attempted

- Degree of polynomial features: Experimentation with degrees  $d = 2, 3, \dots$
- Feature normalization for stability.
- Exploratory Data Analysis to understand data distribution and relationships.

#### **4.4.2 Rejected Approaches**

- High-degree polynomials leading to overfitting.
- Training without normalization causing numerical instability.

### **4.5 Results and Conclusion**

- Polynomial Regression successfully captured non-linear relationships in the data.
- Results demonstrated the importance of selecting an appropriate polynomial degree to balance bias and variance.

## 5 Neural Network Implementation for Binary Labels

### 5.1 Introduction

This section implements a simple feedforward neural network with one hidden layer to perform binary classification tasks.

### 5.2 Objective

The objective of this project is to:

- Implement a feedforward neural network from scratch using Python and NumPy.
- Train the network on a binary classification dataset.
- Evaluate the model using metrics such as binary cross-entropy loss.
- Visualize the training process through cost vs. iterations plots.

### 5.3 Implementation Details

#### 5.3.1 Architecture

The network consists of:

- An input layer with dimensions equal to the number of features.
- A hidden layer with a specified number of neurons and sigmoid activation.
- An output layer with a single neuron for binary classification.

#### 5.3.2 Training Details

- Binary cross-entropy loss function:

$$\mathcal{L}(y, \hat{y}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- Gradient descent to update weights and biases.

### 5.4 Experimentation

#### 5.4.1 Data Preprocessing

- Handled missing values by replacing them with column means.
- Normalized features using StandardScaler.
- Verified data quality by checking for NaN and infinite values.

### 5.4.2 Hyperparameters

- Learning Rate: 0.01
- Hidden Layer Neurons: 4
- Iterations: 1500

## 5.5 Cost Function Convergence

The plot of the cost function vs. iterations is shown below:




Figure 4: Cost function vs. Iterations

## 5.6 Results and Conclusion

- The network successfully learned to classify binary labels with high accuracy.
- The importance of careful data preprocessing was demonstrated.

## 6 Neural Network Implementation for Class Labels

### 6.1 Introduction

Two-layer neural networks are foundational models in deep learning that consist of an input layer, a single hidden layer, and an output layer. This project demonstrates a two-layer neural network's implementation to classify multi-class data.

### 6.2 Objective

The objective of this project is to:

- Implement a two-layer feedforward neural network using Python and NumPy.
- Train the network on a multi-class classification dataset.
- Evaluate the model using cross-entropy loss and accuracy metrics.
- Visualize the training loss over iterations.

### 6.3 Implementation Details

#### 6.3.1 Architecture

The network consists of:

- Input Layer: Processes the feature vectors.
- Hidden Layer: Uses ReLU activation for non-linear transformations.
- Output Layer: Applies softmax activation for multi-class probabilities.

#### 6.3.2 Training Process

- Forward propagation to compute activations.
- Backward propagation to compute gradients for weight updates.
- Optimization using gradient descent.

### 6.4 Experimentation

#### 6.4.1 Data Preprocessing

- Normalized features to zero mean and unit variance.
- One-hot encoded class labels for multi-class targets.

#### 6.4.2 Hyperparameters

- Hidden Layer Size: 128 neurons.
- Learning Rate: 0.01.
- Number of Epochs: 100.

## 6.5 Cost Function Convergence

The plot of the cost function vs. iterations is shown below:




Figure 5: Cost function vs. Iterations

## 6.6 Results and Conclusion

- The network achieved high accuracy on the test dataset, demonstrating its ability to learn complex patterns.
- Visualizations of the loss curve showed steady convergence, validating the training process.